

ACKNOWLEDGMENT

First and foremost, we would like to thank Lord Almighty for making everything possible on time for us and his unseen guidance throughout our work. We should like to thank our parents for their love, support and encouragement.

We express our sincere thanks to **Dr. Roobin V. Varghese**, Principal & Professor, College of Engineering and Management Punnapra and **Dr.Suresh Kumar N**, Associate professor and Head of the department of Computer Science and Engineering, College of Engineering and Management Punnapra, for their kind cooperation and encouragement.

We express our sincere gratitude to our internal guide **Ms. Sneha S**, Assistant professor of the department of Computer Science and Engineering, College of Engineering and Management Punnapra and coordinators **Ms. Smitha M Jasmine**, Assistant Professor, Department of Computer Science and Engineering, College of Engineering and Management Punnapra and **Ms. Remyamol R**, Assistant Professor, Department of Computer Science and Engineering, College of Engineering and Management Punnapra, for their positive attitude, intense support, guidance, encouragement and help throughout our work.

We would like to thank all other faculty members and all our friends for their help and support in carrying out this work successfully.

AMAL S KUMAR
ALEESHA M
AMBILI AMBARAN S
ASWIN S TUSHAR

ABSTRACT

Reconstructing accurate and photorealistic 3D scenes from 2D images is a fundamental challenge in computer vision, particularly when real-time interaction and visual fidelity are both required. Neural Radiance Fields (NeRF) have demonstrated remarkable capability in capturing fine geometric and lighting details through implicit volumetric representations, but their high computational cost results in slow training and rendering, limiting practical deployment. Conversely, 3D Gaussian Splatting represents scenes explicitly using point-based Gaussians, enabling fast rasterization and real-time rendering, though often with increased storage requirements and reduced detail in complex regions. This project introduces a hybrid 3D reconstruction framework that integrates NeRF-based volumetric modeling with Gaussian Splatting-based rendering to leverage the strengths of both methods. The proposed approach aims to produce high-quality, visually consistent 3D reconstructions from unstructured 2D images while achieving real-time rendering performance, efficient data representation, and improved robustness under sparse-view and varying lighting conditions. The outcome of this work contributes toward scalable and interactive 3D reconstruction systems for applications such as virtual reality, digital twins, and immersive visualization.

TABLE OF CONTENTS

Contents	Page No.
ACKNOWLEDGEMENT	i
ABSTRACT	ii
TABLE OF CONTENTS	iii
LIST OF FIGURES	vi
ABBREVIATIONS	vii
CHAPTER 1 INTRODUCTION	
1.1 NEED OF STUDY.....	1
1.2 AIM.....	1
1.3 SCOPE.....	1
1.4 OBJECTIVES.....	2
CHAPTER 2 LITERATURE SURVEY	
2.1 GENERAL.....	3
2.2. LITERATURE REVIEW OF EARLIER WORKS.....	3

CHAPTER 3 FOCUS OF STUDY

3.1 RESEARCH GAP	5
3.2 PROBLEM STATEMENT	5

CHAPTER 4 SYSTEM ANALYSIS

4.1 SYSTEM REQUIREMENT SPECIFICATIONS.....	6
4.1.1 FUNCTIONAL REQUIREMENTS.....	6
4.1.2 NON FUNCTIONAL REQUIREMENTS.....	6
4.1.3 HARDWARE REQUIREMENTS.....	7
4.1.4 SOFTWARE REQUIREMENTS.....	7
4.1.5 CONSTRAINTS.....	8
4.2 PROJECT SCHEDULING.....	8
4.3 COST ESTIMATION.....	8

CHAPTER 5 SYSTEM DESIGN

5.1 SYSTEM DESIGN DOCUMENT.....	10
5.1.1 SYSTEM OVERVIEW.....	10

5.1.2	SYSTEM ARCHITECTURE.....	10
5.1.3	DATA DESIGN.....	11
5.1.4	COMPONENT DESIGN.....	12
5.1.5	ALGORITHM DESIGN.....	12
5.1.6	INTERFACE DESIGN.....	14
5.1.7	REQUIREMENT MATRIX.....	16

CHAPTER 6 IMPLEMENTATION

6.1	HARDWARE IMPLEMENTATION.....	18
6.2	SOFTWARE IMPLEMENTATION.....	19

CHAPTER 7 TESTING

7.1	GENERAL.....	20
7.2	FUNCTIONAL TESTING.....	21
7.3	PERFORMANCE TESTING.....	21
7.4	ACCURACY EVALUATION.....	22
7.5	INTEGRATION TESTING.....	22
7.6	ROBUSTNESS & ERROR HANDLING.....	22
7.7	CONTINUOUS TESTING & IMPROVEMENT.....	22

CHAPTER 8 RESULTS.....	23
-------------------------------	-----------

CHAPTER 9 CONCLUSIONS

9.1 GENERAL.....	25
------------------	----

9.2 FUTURE ENHANCEMENT.....	25
-----------------------------	----

APPENDICES.....	27
------------------------	-----------

REFERENCES.....	29
------------------------	-----------

LIST OF FIGURES

SL. NO	FIGURE NAME	PAGE NO
1	System architecture	11
2	Flow diagram of proposed approach	13
3	Proposed Architecture of framework	14
4	Interface Design	15
5	Output image	24

ABBREVIATIONS

SRS	-	Software Requirement Specification
NeRF	-	Neural Radiance Fields
3DGS	-	3D Gaussian Splatting
GPU	-	Graphics Processing Unit
CUDA	-	Computer Unified DesignArchitecture
FPS	-	Frames Per Second
CLI	-	Command Line Interface
OS	-	Operating System
SfM	-	Structure from Motion
AI	-	Artificial Intelligence
ML	-	Machine Learning
UI	-	User Interface
API	-	Application Programming Interface
CPU	-	Central Processing Unit

CHAPTER 1

INTRODUCTION

The chapter gives a basic overview of the entire project work. It explains how 3D scenes are reconstructed from multiple 2D images using a unified hybrid pipeline that integrates Neural Radiance Fields (NeRF) and 3D Gaussian Splatting (3DGS). Unlike traditional sequential methods, this project focuses on active co-optimization, where NeRF density informs Gaussian pruning to eliminate artifacts. This chapter outlines the need for the study, along with the aim, scope, and objectives aimed at achieving high-fidelity, real-time 3D reconstruction

1.1 NEED OF STUDY

The need for this study arises from the increasing demand for realistic and efficient 3D scene reconstruction in fields such as virtual reality, gaming, and digital heritage. While NeRF provides excellent geometry, it is slow to render; conversely, 3D Gaussian Splatting is fast but often suffers from "floaters" and "blobs" in empty space. Furthermore, existing pipelines often lack stability during the initialization phase and suffer from high memory overhead. Therefore, there is a critical need for a stabilized, hybrid approach that uses volumetric density to guide point-based rendering, ensuring both structural accuracy and computational efficiency.

1.2 AIM

To develop a stabilized, hybrid 3D reconstruction pipeline that converts 2D images into high-fidelity interactive scenes by utilizing a density-guided co-optimization strategy. This combines the volumetric precision of NeRF with the real-time rasterization speed of 3D Gaussian Splatting, while optimizing memory usage through intelligent data caching and downsampling.

1.3 SCOPE

- Convert multiple 2D images into accurate 3D representations using hybrid neural rendering
- Implement a density-guided pruning mechanism where NeRF spatial data is used to actively refine the Gaussian point cloud.
- Optimize computational performance through lazy-loading memory caches and direct ray sampling to eliminate CPU/GPU bottlenecks.
- Ensure pipeline resilience by implementing automated state-saving and "resume from checkpoint" capabilities.

- Provide a real-time interactive dashboard with CUDA-powered rendering for immediate scene visualization at ~30 FPS.
- Support high-resolution input processing through automatic image downsampling to maintain hardware stability.

1.4 OBJECTIVES

- To reconstruct accurate 3D models from 2D image datasets using a stabilized optimization pipeline.
- To utilize NeRF for high-quality volumetric representation and spatial density mapping.
- To apply Aggressive 3D Gaussian Densification and Scale Clamping for sharp, detailed real-time rendering.
- To develop a co-optimization engine that targets and prunes Gaussians in low-density regions (density < 0.01) to eliminate artifacts.
- To optimize the system for interactive performance by implementing chunked evaluation and SH-degree (Spherical Harmonics) warmup strategies.
- To design a user-centric interface for manual parameter control and real-time ETA tracking during the reconstruction process.

CHAPTER 2

LITERATURE SURVEY

2.1 GENERAL

The 3D Reconstruction from 2D Photos using Neural Radiance Fields (NeRF) and Gaussian Splatting involves an extensive review of existing research papers, publications, and studies related to 3D reconstruction, computer vision, and neural rendering techniques. The survey aims to understand the different methods and technologies used for converting 2D images into realistic 3D scenes, analyze their advantages and limitations, and identify possible improvements for developing an efficient and high-quality 3D reconstruction system.

2.2. LITERATURE REVIEW OF EARLIER WORKS

The field of 3D scene reconstruction has undergone a transformative journey, transitioning from discrete geometric mapping to continuous neural functions. Early methodologies relied heavily on **Structure-from-Motion (SfM)** [3] as detailed by **Schönberger and Frahm (2016)** to construct sparse point clouds. These systems typically utilized the **Scale-Invariant Feature Transform (SIFT)** [9] developed by **David Lowe (2004)** to match keypoints across various 2D frames. While these traditional approaches provided mathematically accurate geometry, they fundamentally lacked the photorealism required for truly immersive reality. This gap led to the introduction of **Neural Radiance Fields (NeRF)** [1] by **Mildenhall et al. (2020)**, which redefined the field by using a Multi-Layer Perceptron (MLP) to represent scenes as continuous volumetric functions.

The evolution of NeRF led to significant architectural refinements. **Barron et al. (2021)** introduced **Mip-NeRF** [18] to handle multiscale representations and anti-aliasing, while their subsequent work on **Mip-NeRF 360** [5] (2022) addressed the challenges of unbounded scenes. To combat the slow rendering speeds of early models, **Müller et al. (2022)** developed **Instant-NGP** [4], utilizing multiresolution hash encoding to reduce training times drastically. Other optimization strategies emerged, such as **PlenOctrees** [17] by **Yu et al. (2021)** and **BakedSDF** [12] by **Hedman et al. (2023)**, which focused on meshing and caching neural fields for real-time synthesis. In the 2025 research cycle, **Tianqi Ding (2025)** furthered this efficiency by demonstrating that coreset-driven neural pruning can reduce model size by 50% without compromising reconstruction accuracy.

As the research momentum shifted toward real-time accessibility and explicit geometry, **3D Gaussian Splatting (3DGS)** [2] by **Kerbl et al. (2023)** emerged as a dominant alternative, moving away from implicit volumes toward explicit Gaussian primitives. This transition has enabled new domains of scene manipulation, as categorized by **Shuting He (2025)** in segmentation and editing. To facilitate this, **Vickie Ye (2025)** developed "gsplat," an optimized CUDA back-end for **PyTorch** [6] (Paszke et al., 2019). Recent hybrid approaches have sought to align these Gaussians with surfaces for better reconstruction; **Guédon and Lepetit (2023)** introduced **SuGaR** [7] for surface-aligned splatting, while **Qi-Yuan Feng (2025)** developed "EVSplitting" to ensure visual consistency during complex geometric edits like slicing and engraving.

The challenge of sparse-view reconstruction—synthesizing scenes from limited data—has seen significant breakthroughs. While **Zhang et al. (2020)** improved few-shot synthesis with **NeRF++ [13]**, more recent 2025 models like "**DT-NeRF**" by **Bo Liu (2025)** integrate Diffusion Models and Transformers to maintain structural integrity. **Haitian Liu (2025)** also addressed these constraints by mitigating scale ambiguity in sparse-view 3DGS frameworks. For complex materials and reflections, **Albert Barreiro (2025)** performed comparative analyses on specular surfaces, while **M. Maharani (2025)** evaluated the trade-offs between traditional **Multi-View Stereo (MVS)** and neural paradigms.

Modern applications of these technologies have expanded into dynamic environments and industrial automation. **Yanjun Ma (2025)** and **Ma et al. (2025)** have optimized visual SLAM systems using NeRF-based background mapping, while **Chaoran Feng (2025)** introduced "**AE-NeRF**" for robust event-based reconstruction in noisy conditions. **Xunzhi Zheng (2025)** furthered this by developing "**Flow-NeRF**," which jointly learns geometry and optical flow to resolve ambiguities in camera pose estimation. In the industrial sector, **Eyob T. Mengiste (2025)** and **Zeyu Li (2025)** have successfully applied these neural rendering methods to construction site monitoring and building curtain wall maintenance.

Our project, **Gausfer 3D**, draws technical inspiration from the grid-based sampling methods of **Plenoxels [8] (Fridovich-Keil et al., 2022)** and the object-level reasoning of **PointNet [15] (Qi et al., 2017)**. By integrating a 3D-aware positional modeling mechanism, similar to the geometry-optimized strategies of **Zechuan Li (2024)**, **Gausfer 3D** aims to resolve the misalignments and "floaters" common in standard splatting. We utilize the **Structural Similarity Index (SSIM) [10]** by **Wang et al. (2004)** to ensure our hybrid model achieves superior visual fidelity. By bridging the volumetric consistency of NeRF with the real-time speed of 3DGS, **Gausfer 3D** provides a robust, hardware-efficient solution for high-quality 3D digitization.

CHAPTER 3

FOCUS OF STUDY

The focus of this study is the development of a stabilized hybrid 3D reconstruction system that transforms unstructured 2D photographs into realistic, interactive 3D environments. The study centers on the active co-optimization of Neural Radiance Fields (NeRF) and 3D Gaussian Splatting (3DGS). By utilizing NeRF's volumetric density to guide Gaussian refinement, the project aims to eliminate common neural rendering artifacts—such as "floaters" and "blobs"—while maintaining a real-time rendering performance of ~30 FPS

3.1 RESEARCH GAP

Existing 3D reconstruction methods face a significant trade-off between visual fidelity and rendering efficiency. While NeRF produces high-quality novel view synthesis, its slow rendering speed limits interactivity. Conversely, 3D Gaussian Splatting (3DGS) offers real-time speeds but often suffers from geometric instability, leading to "floaters" in empty space and a loss of fine detail in complex shapes.

A critical gap identified in current academic frameworks is pipeline reliability and resource management. Many existing implementations are prone to:

Initialization Failures: Crashing during the sparse-to-dense point cloud transition (e.g., "no dense points found" errors).

Memory Bottlenecks: High VRAM usage and CPU stalls during ray-marching and high-resolution (4K) processing.

Lack of State Preservation: The inability to resume training after a system interruption, leading to lost computational progress.

There is a clear need for a framework that not only merges NeRF and 3DGS but also introduces density-aware pruning and memory-optimized training to make the process resilient and hardware-efficient

3.2 PROBLEM STATEMENT

This project proposes a hybrid 3D reconstruction framework that bridges the photorealistic accuracy of NeRF with the real-time rendering speed of 3D Gaussian Splatting. It efficiently transforms ordinary 2D images into high-fidelity, interactive 3D environments.

CHAPTER 4

SYSTEM ANALYSIS

System analysis is the process of studying a proposed system in detail to understand its objectives, requirements, components, constraints, and overall feasibility. It helps in identifying what the system should do, how it should perform, and what resources are needed for implementation. In the proposed project, system analysis is carried out to examine the requirements of the hybrid 3D reconstruction framework that converts multiple 2D photographs into a navigable 3D scene using Neural Radiance Fields and 3D Gaussian Splatting.

4.1 SYSTEM REQUIREMENT SPECIFICATIONS

The Software Requirement Specification outlines the functional and non-functional requirements necessary for the proposed system to operate effectively. The system is designed to transform multiple 2D images into a high-fidelity 3D scene by integrating NeRF and 3D Gaussian Splatting while supporting interactive visualization and efficient rendering..

4.1.1 FUNCTIONAL REQUIREMENTS

- **FR1: Upload Multi-View Images**
The system shall allow the user to upload multiple 2D images of a scene captured from different viewpoints. The accepted formats include JPEG and PNG. After successful upload, the system shall provide confirmation to the user.
- **FR2: Camera Pose Estimation**
The system shall estimate camera intrinsic and extrinsic parameters from the uploaded images. This step is necessary for accurate spatial understanding of the scene and is performed through preprocessing and Structure-from-Motion techniques.
- **FR3: Generate 3D Model using NeRF**
The system shall generate a volumetric 3D scene representation using Neural Radiance Fields. It uses the input images and estimated camera poses to learn scene geometry and appearance.
- **FR4: Interactive Scene Navigation**
The system shall allow users to interact with the reconstructed 3D scene through zooming, rotating, and panning. This enables real-time viewing and exploration of the generated scene.
- **FR5: Save Reconstructed Scene**
The system shall allow the user to save or export the reconstructed 3D scene for future use. The exported output may be in a suitable model format such as .ply.

4.1.2 NON-FUNCTIONAL REQUIREMENTS

- **Performance Requirements**
The system shall reduce training time compared to traditional NeRF-only methods by using a hybrid NeRF–Gaussian Splatting approach. It shall support real-time or near real-time inference

and maintain stable rendering performance for varying scene complexities.

➤ **Safety Requirements**

The system shall avoid abrupt crashes caused by memory overflow or excessive GPU utilization. Proper exception handling shall be implemented to ensure safe execution.

➤ **Security Requirements**

The system shall process all image data and reconstructed outputs locally without external transmission. It shall also protect project data from unauthorized access, modification, or deletion.

➤ **Reliability**

The system shall provide stable and consistent 3D reconstruction results. It shall operate without unexpected failures during prolonged training sessions.

➤ **Usability**

The system shall provide a simple and intuitive workflow for image upload, training, rendering, and visualization. The generated output shall be clearly viewable by the user.

➤ **Maintainability**

The code shall follow a modular and well-documented structure so that updates to one component do not affect the entire system.

➤ **Scalability**

The system shall support increasing numbers of input images and different scene sizes with minimal performance degradation.

➤ **Portability**

The system shall be executable on multiple operating systems such as Windows and Linux with minimal changes, provided Python and GPU support are available.

4.1.3 HARDWARE REQUIREMENTS

The proposed system requires a GPU-enabled environment because both neural reconstruction and Gaussian rendering involve intensive computation. The hardware requirements are:

1. NVIDIA CUDA-supported GPU
2. Multi-core processor
3. Minimum 8 GB RAM or above
4. Adequate secondary storage for datasets and output files
5. Display monitor, keyboard, and mouse for user interaction

4.1.4 SOFTWARE REQUIREMENTS

The software requirements for the proposed system are:

1. Operating System – Windows or Linux
2. Programming Language – Python
3. Deep Learning Framework – PyTorch

4. GPU Acceleration Platform – CUDA
5. Structure-from-Motion Tool – COLMAP
6. NeRF Frameworks
7. Gaussian Splatting libraries
8. Visualization and rendering tools

4.1.5 CONSTRAINTS

The proposed system has certain limitations and constraints. The performance of the system depends heavily on the availability of CUDA-supported GPU hardware. Reconstruction quality is strongly influenced by the quality, overlap, and angular diversity of the input images. Sparse-view inputs may reduce the accuracy of the reconstructed model. The current system is mainly intended for static scenes, and dynamic scene support is not included in the present scope. In addition, large and high-resolution scenes may increase memory usage and processing time.

4.2 PROJECT SCHEDULING

Project scheduling helps in organizing the work into different phases and completing them in a systematic way. For the proposed 3D reconstruction system, the project can be divided into the following stages: literature survey, requirement analysis, system design, preprocessing and camera pose estimation, NeRF-based reconstruction, Gaussian Splatting integration, interface development, testing, and documentation. Proper scheduling ensures timely completion of the project and balanced use of available resources.

4.3 COST ESTIMATION

Cost estimation gives an approximate idea about the resources needed for implementing the project. Since the proposed system is mainly intended for academic and research purposes, the major cost components are related to hardware, software environment, storage, and development effort

4.3.1 HARDWARE COST:

The major hardware cost is associated with a GPU-enabled system. Since neural rendering and rasterization require high computational power, a CUDA-supported NVIDIA GPU is necessary. Additional cost may include RAM upgrades and storage devices for handling large image datasets and output models.

4.3.2 SOFTWARE DEVELOPMENT COSTS:

Most of the software frameworks used in the project, such as Python, PyTorch, COLMAP, and many research libraries, are open source. Therefore, software licensing cost is minimal. However, development requires time, technical expertise, and computational resources.

4.3.3 SYSTEM INTEGRATION COSTS:

System integration cost includes combining preprocessing modules, camera pose estimation tools, reconstruction modules, and rendering components into a unified workflow. Testing and compatibility management also contribute to this cost.

4.3.4 MAINTENANCE AND SUPPORT COSTS:

Maintenance cost includes updating libraries, improving modules, debugging, and managing compatibility with future versions of frameworks and drivers

4.3.5 MISCELLANEOUS COSTS:

Miscellaneous costs may include internet access for downloading frameworks, backup storage devices, electricity charges for prolonged GPU usage, and contingency resources for experimentation and testing. Since the system is developed under academic settings using open-source tools, the overall cost is moderate

CHAPTER 5

SYSTEM DESIGN

5.1 SYSTEM DESIGN DOCUMENT

The System Design Document for the Gausfer 3D project outlines the architectural and functional aspects of the system in detail. This serves as a comprehensive guide for developing the System Model, ensuring alignment with the project objectives of high-fidelity 3D reconstruction and real-time rendering

5.1.1 SYSTEM OVERVIEW

Gausfer 3D is a hybrid 3D reconstruction pipeline that bridges Neural Radiance Fields (NeRF) and 3D Gaussian Splatting (3DGS). It transforms simple 2D photographs into navigable, immersive 3D environments by combining the volumetric consistency of NeRF with the interactive speed of Gaussian primitives.

5.1.2 SYSTEM ARCHITECTURE

The system architecture ensures efficient scene reconstruction, rapid geometric convergence, and seamless real-time rendering, delivering a robust model for spatial computing

α (Alpha): Opacity limit used to prune transparent artifacts.

$\mathbf{R}$: 3X3 rotation matrices converted from COLMAP quaternions

T_{grad} : Threshold for positional gradients used in densification

5.1.4 COMPONENT DESIGN

The system is divided into modular components: a 5-layer MLP for neural scaffolding, a NeRF-to-Gaussian Bridge for seeding, and a CUDA-accelerated rasterization kernel for output.

5.1.5 ALGORITHM DESIGN

The core "Guided Densification Logic" identifies high-gradient Gaussians and validates them against the NeRF density field.

Step 1: Disable gradient tracking

Step 2: $\mathbf{g} \leftarrow \nabla \mathbf{x}_{\text{gaussians}}$

Step 3: $M_{\text{high}} \leftarrow \|\mathbf{g}\|_2 > \tau$

Step 4: if $\exists M_{\text{high}}$ then

Step 5: $\mathbf{x}_{\text{prob}} \leftarrow \mathbf{x}_{\text{gaussians}}[M_{\text{high}}]$

Step 6: $\rho \leftarrow \text{NeRF}(\mathbf{x}_{\text{prob}})$

Step 7: $M_{\text{valid}} \leftarrow \rho > \delta$

Step 8: if $\exists M_{\text{valid}}$ then

Step 9: Construct combined mask M_{combined}

Step 10: Perform Gaussian densification and splitting

Step 11: end if

Step 12: end if




Fig: Flow Diagram of the proposed approach




Fig: Use Case Diagram

Fig: Proposed Architecture of our framework

5.1.3 INTERFACE DESIGN

The interface focuses on usability by automating complex tasks. Users provide consumer-grade video imagery, and the system handles pose estimation and reconstruction automatically, providing a real-time navigable view of the scene.

	baseline volumetric density field for geometric guidance.	5-layer MLP.	geometric convergence with a PSNR of 32.5 dB .
REQ-04	Hybrid Densification: Gaussians must only densify in confirmed solid surface areas.	Monitor densification in empty space regions.	Cross-references high-gradient Gaussians (∇_x) with the NeRF scaffold ($\sigma_{\text{NeRF}} > 5.0$) before splitting.
REQ-05	Artifact (Floater) Pruning: Primitives in empty space or with low opacity must be removed.	Observe model cleanup in unbounded or noisy scene regions	Aggressively prunes primitives if opacity $\alpha < 0.05$ or NeRF identifies space as empty ($\sigma_{\text{NeRF}} < 1.0$)
REQ-06	Real-Time Rendering: System must support interactive novel view synthesis	Benchmark rendering speed on a consumer-grade GPU.	Custom CUDA kernels for rasterization enable real-time performance of >100 FPS.
REQ-07	Memory Optimization: The system must operate within a fixed VRAM limit	Monitor GPU memory usage during 5000-step co-optimization.	Employs ray-chunking (4096) and strict density thresholds to maintain an 8 GB VRAM budget.

CHAPTER 6

IMPLEMENTATION

The implementation of Spectrafusion 3D focuses on a synergistic integration of implicit neural representations and explicit geometric primitives. The framework is designed for maximum resilience, utilizing automated checkpointing and optimized memory repositories to ensure high-fidelity reconstruction on consumer-grade hardware.

6.1 HARDWARE IMPLEMENTATION

The hardware implementation is strategically designed to democratize 3D reconstruction by making it functional on standard computing resources. The primary computational engine is a CUDA-enabled NVIDIA GPU, utilized for parallelizing the intensive requirements of both NeRF training and Gaussian rasterization.

A critical challenge addressed in this implementation is the management of VRAM during high-resolution processing. To resolve this, the system incorporates Automatic Image Downsampling, which resizes all input datasets to a maximum dimension of 1200px in [dataset_loader.py](#). This reduces the computational payload on the rasterizer by approximately 10x for 4K inputs, allowing the system to operate comfortably within an 8 GB VRAM budget.

To further optimize throughput, the hardware pipeline utilizes a Lazy-loading Memory Cache within the [get_training_batch](#) module. This implementation pulls training elements rapidly from a CPU RAM repository, which funnels matrix calculations into the core processors and prevents GPU "starvation" or driver stalls. Furthermore, the hardware supports a CUDA-powered Real-Time Viewer capable of rendering at ~30 FPS with full orbit and zoom controls, providing an immediate feedback loop for the user during the optimization process.

6.2 SOFTWARE IMPLEMENTATION

The software implementation is a multi-stage pipeline developed using PyTorch and Streamlit. It has been significantly updated to include "Training Resilience" and "Density-Guided Sharpness," moving through three primary phases:

6.2.1 Data Acquisition and Stability

The process begins with the extraction of frames from raw video, which are then processed by a custom COLMAP wrapper for camera pose estimation. A key software innovation in this phase is the stabilization of the Gaussian Seeding logic. We resolved common initialization crashes by re-engineering [bridge_converter.py](#) to handle sparse-to-dense point cloud transitions more robustly, ensuring that "no dense points found" errors no longer terminate the pipeline.

Additionally, an Automated State Saving mechanism was implemented, allowing the software to generate checkpoints that users can "Resume" from if the process is interrupted.

6.2.2 Neural Scaffolding and Ray Optimization

The software utilizes a 5-layer Multi-Layer Perceptron (MLP) to establish a baseline density field. To remove a massive CPU bottleneck found in standard models, we replaced the full-meshgrid ray calculation in `nerf_trainer.py` with Direct Ray Sampling. This allows the software to query the NeRF density (σ) much faster. This density field acts as a "geometric guide," identifying regions where the probability of matter is high ($\sigma > 15.0$) to seed the initial Gaussian primitives.

6.2.3 Hybrid Co-Optimization and Detail Warmup

The final stage is a 5000-step hybrid optimization loop governed by two critical algorithms:

- **Aggressive Densification:** The system now triggers densification every 25 steps (twice as frequently as previous versions). This helps fill complex shapes like furniture and textured surfaces much faster.
- **SH-Degree Warmup:** To prevent "floaters" and messy artifacts, the software implements a warmup strategy for Spherical Harmonics. It initially focuses on solid colors and geometry before gradually introducing lighting and reflection details.
- **Density-Guided Pruning:** Primitives are cross-referenced against the NeRF scaffold. If a Gaussian exists in a region where NeRF density is $\sigma < 0.01$, it is immediately pruned. Coupled with Tighter Scale Clamping, this ensures that the final model is sharp and free of "blobs," resulting in high-fidelity reconstructions of intricate details like carpets, slippers, and electronic equipment.

CHAPTER 7

TESTING

7.1 GENERAL

Testing is a critical phase in the development of the Spectrafusion 3D framework. It ensures the hybrid reconstruction pipeline is accurate, stable, and resilient to hardware constraints. This process involves validating the transition from 2D images to 3D representations, specifically focusing on the new density-guided pruning and state-preservation features.

7.2 FUNCTIONAL TESTING:

This testing verifies that each core functionality of the Gausfer 3D system operates according to the specified functional requirements:

- **FR1: Dataset Optimization:** Verified that `dataset_loader.py` correctly resizes images to 1200px and that the Lazy-loading RAM cache effectively funnels data without GPU stalls.
- **FR2: Stable Camera Estimation:** Tested the integration with COLMAP and the stabilized `bridge_converter.py` to ensure that Gaussian seeding no longer fails with "no dense points found" errors.
- **FR3: Hybrid Modeling:** Validated that the NeRF-guided pruning engine correctly identifies and removes "floaters" in regions where density is $\sigma < 0.01$.
- **FR4: Real-Time Navigation:** Tested the CUDA-powered Orbit Viewer within the Streamlit dashboard to ensure a consistent 30 FPS experience.
- **FR5: Resilience & Checkpointing:** Confirmed that the system creates automatic state saves and that the "Resume from Checkpoint" toggle successfully recovers training after an interruption.

7.3 PERFORMANCE TESTING:

Performance testing evaluates the system's efficiency and hardware utilization:

- **Rendering Speed:** Benchmarked the rasterization process using `quick_render.py` to ensure it maintains interactive speeds despite aggressive densification.
- **Training Efficiency:** Measured the impact of Direct Ray Sampling in `nerf_trainer.py`, confirming a significant reduction in CPU overhead and faster epoch completion.
- **Memory Load:** Monitored VRAM usage during the 5000-step co-optimization, ensuring that Chunked NeRF Evaluation prevents Out-of-Memory (OOM) crashes even with high Gaussian counts..

7.4 ACCURACY EVALUATION:

This phase assesses the visual fidelity using specific calibration scripts:

- **Calibration Verification:** Used `get_bbox.py` to verify that the Dynamic Bounding Box correctly encapsulates the target scene without cutting off geometry.
- **Geometric Sharpness:** Specifically tested the Tighter Scale Clamping logic to ensure that complex objects (e.g., table edges, slippers) appear sharp and intentional rather than "blobby."
- **Detail Warmup:** Verified the SH-Degree Warmup strategy, confirming that reflection artifacts are minimized by focusing on basic geometry during the early training steps.

7.5 INTEGRATION TESTING:

Integration testing focuses on the coordination between various system modules and external libraries:

- **Optimized Training Loop:** Verified the integration of the `update_optimizer` routine with the UI thread, ensuring that array resizing for existing points occurs without dropping structural gradients.
- **NeRF-to-Gaussian Bridge:** Tested the data transfer integrity between the 5-layer MLP and the Gaussian primitives during the every-25-steps densification cycle.

7.6 ROBUSTNESS & ERROR HANDLING:

Robustness testing assesses the system's resilience against non-ideal conditions:

- **Pipeline Stability:** Performed "stress tests" by interrupting the training loop to verify that the automated checkpointing allows for a seamless resume.
- **Initialization Resilience:** Verified that the system provides informative error messages and recovery options instead of crashing when encountering sparse input data during the seeding phase.

7.7 CONTINUOUS TESTING AND IMPROVEMENT:

The following specialized scripts were implemented for continuous validation:

- **`quick_render.py`:** For rapid verification of current scene geometry without full optimization.
- **`test_hybrid_optim.py`:** To benchmark the interaction between NeRF density checks and Gaussian pruning accuracy.

CHAPTER 9

CONCLUSION

9.1 GENERAL

The Spectrafusion 3D project successfully establishes a resilient hybrid framework for 3D reconstruction by synthesizing the volumetric accuracy of Neural Radiance Fields (NeRF) with the real-time efficiency of 3D Gaussian Splatting (3DGS). The project achieves its primary objective: overcoming the "floater" and "blob" artifacts typical of standalone splatting methods through the implementation of a density-guided co-optimization engine.

The system effectively fulfills all specified functional requirements. By integrating Direct Ray Sampling and Lazy-loading Memory Caching, the pipeline demonstrates that high-fidelity 3D reconstruction is achievable on consumer-grade hardware within an 8 GB VRAM budget. The inclusion of Automated State Saving and a Streamlit-based Dashboard transforms the research into a stable, user-centric application, allowing for session recovery and real-time parameter tuning.

Ultimately, the project delivers a balanced trade-off between geometric precision and interactive performance, achieving stable rendering at ~30 FPS. This framework serves as a robust platform for spatial computing research, proving that hybrid neural rendering is the key to creating immersive, artifact-free digital twins from ordinary 2D photographs.

9.2 FUTURE ENHANCEMENT

While the current system provides a stabilized and high-performance pipeline, the following areas are identified for future research and expansion:

- **Dynamic 4D Reconstruction:** The current system is optimized for static environments. Future work will explore Temporal NeRF techniques and dynamic Gaussian deformation to reconstruct scenes with moving objects or changing lighting in real-time
- **Generative Geometry Infilling:** To handle sparse-view datasets where image overlap is minimal, Generative AI and Diffusion-based priors can be integrated to "hallucinate" missing geometry and textures, ensuring a complete 3D model even with limited input
- **Advanced Primitive Compression:** Although the current pipeline uses scale clamping to manage density, future enhancements could include Vector Quantization or Entropy Coding to compress the millions of Gaussian primitives, reducing the storage footprint for large-scale urban environments.
- **Semantic Scene Understanding:** Integrating Semantic Segmentation into the hybrid loop would allow the system to recognize specific objects (e.g., "chair," "wall"). This would enable selective densification and independent manipulation of objects within the 3D scene.
- **Cloud-Edge Collaborative Processing:** To support mobile devices, the system could be extended to a distributed architecture where heavy NeRF scaffolding is performed on cloud servers while real-time Gaussian rendering is handled locally on the edge device.

- **Material and Physics Integration:** Future iterations could move beyond color and density to estimate PBR (Physically Based Rendering) materials. This would allow for realistic re-lighting of the reconstructed scene and physics-based interactions with the 3D model in virtual environments.

The proposed future enhancements will improve the overall performance, efficiency, and usability of the system. By addressing current limitations such as static scene support, storage requirements, and dependency on input data, the system can be made more robust and scalable. Incorporating advanced techniques and user-friendly features will further enhance reconstruction quality and expand its applications in real-world scenarios.

APPENDICES

The testing phase ensures that the transition from 2D image datasets to an interactive 3D scene is mathematically sound and computationally efficient. Testing covers the full pipeline, from initial feature matching in COLMAP to the final real-time rendering of the Gaussian splats.

1. Technical Details

The Gausfer 3D framework is implemented using a hybrid architecture that integrates implicit neural representations with explicit volumetric rasterization.

- Development Environment: Python 3.9+, PyTorch 2.0+ with CUDA 11.8 support.
- Core Libraries: differential-gaussian-rasterizer by graphdeco-inria: For high-performance Gaussian rasterization.
- COLMAP: Used as the backend for Structure-from-Motion (SfM) to generate initial camera poses and sparse point clouds.
- Nerfacc / PyTorch Lightning: For efficient volumetric sampling and training pipeline management.
- Data Structures: 3D Gaussians are parameterized by position (mean), covariance (scale and rotation), opacity, and color (Spherical Harmonics).
- Hardware Setup: NVIDIA RTX 40 Series GPU (8GB+ VRAM recommended), 16GB System RAM, and Multi-core CPU for parallel image processing.

2. Code Snippets

The project utilizes a modular and well-documented code structure to ensure maintainability and scalability. Key implementation areas include:

- Automated frame extraction and filtering.
- Neural network training using PyTorch.
- High-performance rasterization using CUDA kernels

Example Code Snippets:

```
# Conceptual example of the Density Validation & Pruning logic
if positional_gradient > threshold:
    if nerf_density > 15.0:
        # High gradient and matter present: Split Gaussian to add detail
        gaussian.split()
    elif nerf_density < 0.05:
        # Space is empty according to NeRF: Prune/Delete Gaussian
        gaussian.delete()
```

3. Supporting Materials

These materials provide the necessary data and performance evidence for the reconstruction

framework.

- Dataset Specifications: Unstructured multi-view 2D images used as input for training.
- Validation Metrics: Performance is measured using PSNR (Peak Signal-to-Noise Ratio), SSIM (Structural Similarity Index), and LPIPS (Learned Perceptual Image Patch Similarity).
- Visualization Results: Real-time 3D navigation windows and log outputs captured during processing.

Example Supporting Materials:

- PSNR: 32.5 dB (indicates high signal-to-noise ratio and image cleanliness).
- SSIM: 0.89 (shows structural similarity to the ground truth).
- LPIPS: 0.18 (low score means images look natural to the eye).
- Hardware Log: Confirmation that training fits within an 8 GB VRAM budget on consumer-grade GPUs.

4. External Libraries and Modules

The system relies on specialized frameworks to handle the complex computations of neural rendering and 3D geometry.

- PyTorch: The primary framework for neural network training and optimization.
- CUDA: Used for hardware-accelerated parallel computation and efficient rasterization.
- COLMAP: Essential for Structure-from-Motion (SfM) to estimate camera poses

Example External Libraries

- differential-gaussian-rasterizer: A comprehensive open-source library with optimized CUDA kernels designed to accelerate the training of 3D Gaussian Splatting (3DGS) methods.
- Torch-NeRF: Utilized for its hash encoding techniques to reduce NeRF's training times.
- SAM/CLIP: Mentioned as 2D foundation models that can facilitate semantic understanding within 3DGS frameworks

REFERENCES

- [1] Vickie Ye in 2025, “gsplat: An open-source library for accelerating 3D Gaussian Splatting and real-time rendering”.
- [2] Shuting He in 2025, “Applications of 3D Gaussian Splatting in segmentation, editing and generation tasks”.
- [3] Haitian Liu in 2025, “Sparse-view 3D reconstruction using 3D Gaussian Splatting techniques”.
- [4] Jiaxin Zhang in 2025, “COB-GS: Boundary-aware 3D Gaussian Splatting for improved scene segmentation”.
- [5] Qi-Yuan Feng in 2025, “EVSplitting: Efficient Gaussian splitting method for 3D scene editing and manipulation”.
- [6] Kyle (Yilin) Gao in 2025, “A comprehensive survey on Neural Radiance Fields from 2020 to 2025”.
- [7] Tianqi Ding in 2025, “Neural pruning techniques to optimize NeRF training and reduce computational cost”.
- [8] Bo Liu in 2025, “DT-NeRF: Integration of diffusion models and transformers for enhanced 3D reconstruction”.
- [9] Yanjun Ma in 2025, “NeRF-based visual SLAM system for dynamic environments”.
- [10] Chaoran Feng in 2025, “AE-NeRF: Event-based Neural Radiance Field for complex scene reconstruction”.
- [11] Xunzhi Zheng in 2025, “Flow-NeRF: Joint learning of geometry, pose and optical flow in neural rendering”.
- [12] Shuai Liu in 2025, “Evolution of 3D reconstruction from traditional methods to NeRF and Gaussian Splatting”.
- [13] Albert Barreiro in 2025, “Comparative study of Ref-NeRF and NRFF for reflective surface modeling”.
- [14] Kibon Ku in 2025, “SC-NeRF for plant phenotyping and agricultural 3D reconstruction”.
- [15] Zechuan Li in 2024, “GO-N3RDet: Geometry optimized NeRF for 3D object detection”.
- [16] M. Maharani in 2025, “Comparison of Multi-View Stereo and NeRF using PCA analysis”.
- [17] Chen Tao in 2025, “Overview and applications of Neural Radiance Fields in 3D reconstruction”.

- [18] Zeyu Li in 2025, “NeRF-based 3D reconstruction using UAV imagery for construction analysis”.
- [19] Zhihong Chen in 2025, “Hybrid encoder-based NeRF for large-scale sparse-view reconstruction”
- [20] Eyob T. Mengiste in 2025, “Evaluation of NeRF and 3D Gaussian Splatting for texture and surface modeling”.